

Structural Entropy: A Decade Survey of Theory, Algorithms, Learning Paradigms, and Applications

TGINA Journal Manuscript (V1.1)

Dingli Su^{1*} Hao Peng¹ Yicheng Pan^{1,†} Angsheng Li¹

¹Beihang University, Beijing, China

This article is an expanded version of the conference paper “A Survey of Structural Entropy: Theory, Methods, and Applications” published at the 34th International Joint Conference on Artificial Intelligence (IJCAI 2025).

Draft line: **V1.1 traceable writing-model rewrite for comparison with V0.0 and V1.0**

Abstract

Graphs and networks are now a common language for artificial intelligence, from social and biological interaction systems to knowledge graphs, retrieval pipelines, and autonomous-agent workflows. A central challenge in these settings is to identify meaningful organization across scales rather than only local interactions. *Structural Entropy* (SE) addresses this challenge by extending information-theoretic reasoning from pairwise uncertainty to hierarchical structure: it measures how efficiently a graph can be encoded by a multilevel partition and, through minimization, reveals candidate structural abstractions of the system.

Over the past decade, SE has developed from a community-detection principle into a broader framework spanning graph clustering, graph pooling, graph structure learning, reinforcement learning, bioinformatics, social-network analysis, and pattern recognition [16, 66, 91, 97, 41, 101]. In this survey we review that development. We summarize the mathematical foundations of SE, organize the literature into major methodological families, compare SE with neighboring graph objectives such as modularity and flow-based compression, and examine how the field is evaluated in practice. We also discuss open problems in scalability, differentiable optimization, heterogeneous structures, and integration with modern graph-enhanced AI systems. A companion open-source library, `glass-jax`, accompanies this survey as a reproducibility artifact.

Keywords: structural entropy, graph clustering, hierarchical community detection, graph neural networks, reinforcement learning, information theory.

*Equal contribution.

†Corresponding author: yichengp@buaa.edu.cn

1 Introduction

Information is a basic organizing principle of artificial intelligence. Learning systems compress observations into representations, decision systems reduce uncertainty about actions, and graph-based systems try to expose structure that is not immediately visible from local relations alone. Much of classical information theory, however, is local in flavour: it quantifies uncertainty in symbols, variables, or pairwise dependencies. **Structural Entropy (SE)** [16] asks a different question. Instead of only measuring uncertainty at the level of isolated variables, it studies the uncertainty embedded in the organization of a whole system.

This perspective is especially natural for graphs. Real graph data are rarely just collections of edges. Social networks contain groups and subgroups, biological systems organize into functional modules, and modern AI pipelines increasingly generate interaction graphs whose usefulness depends on whether they admit meaningful hierarchies. In such settings, one often wants more than a flat partition or a predictive score. One wants a structural description of the system: which components belong together, at what scale, and according to what coding principle.

SE provides one answer to this problem. In its original form, it connects graph organization to coding trees and interprets good hierarchical partitions as efficient information-theoretic descriptions of the graph [16]. It is both a measure of residual structural uncertainty and an optimization principle for discovering efficient hierarchical decompositions. That combination has made SE useful not only for community detection, where it first gained attention, but also for graph pooling, augmentation, clustering, control, and domain-specific discovery tasks [66, 65, 91, 97, 41, 101].

This perspective is timely for graph intelligence because contemporary systems increasingly need more than accurate predictions. They need structural abstractions that are interpretable, reusable, and stable enough to support retrieval, reasoning, and decision making. The need is especially visible in graph-enhanced large language model pipelines, knowledge-graph reasoning systems, tool-use memory, and autonomous-agent interaction graphs. These systems create relational structures whose usefulness depends on whether meaningful hierarchies can be separated from incidental links.

The field has grown quickly since the foundational work of Li and Pan in 2016. What began as a graph-theoretic and coding-theoretic framework for uncovering community structure now includes variants for directed graphs, dynamic settings, differentiable deep models, and decision-making systems [16, 38, 53, 91]. SE-based ideas have appeared in graph neural networks, reinforcement learning, social-bot analysis, genomic data analysis, traffic modelling, and speech or vision applications [72, 4, 103, 61]. At the same time, the literature has become fragmented. Different papers use SE as a discrete objective, a regularizer, a hierarchy prior, a robustness signal, or a task-specific structural score. As a result, readers can now find many applications of the idea, but it is harder to see the shape of the field as a whole.

This survey reviews the first decade of SE research and organizes it around a small number of recurring questions. What exactly does SE measure, and how does it differ from related graph objectives such as modularity, the map equation, spectral criteria, or other entropy-based quantities? What algorithmic families have been developed to minimize or approximate it? How has the idea been adapted to differentiable graph learning and

reinforcement learning? Which application areas have produced the most convincing evidence of value, and how should those claims be evaluated? By bringing these threads together, we aim to provide both a map of the literature and a clearer understanding of what SE contributes as an information-theoretic view of graph intelligence.

Our perspective in this journal version is shaped by two developments. First, SE research is now large enough to support a genuine survey rather than a short position piece. Second, graph intelligence has changed substantially in the same period. Contemporary systems increasingly combine graph structure with foundation models, retrieval pipelines, agent memories, and multimodal interaction logs. These systems create new opportunities for structural abstraction, but they also raise harder questions about interpretability, stability, and evaluation. SE is timely in this setting because it offers a language for reasoning about organization itself, not only prediction.

The present manuscript substantially extends our earlier conference survey [51]. Relative to that version, the journal manuscript expands the theoretical background, broadens the algorithm catalogue, updates the application review through 2026, adds a reproducibility benchmark with the `glass-jax` companion library, and develops a fuller discussion of open problems. The aim is not only to collect papers, but to organize a field that now spans classical graph theory, modern graph learning, and emerging graph-enhanced AI systems.

The remainder of the survey follows that logic. Section 2 positions SE relative to neighboring graph objectives and learning paradigms. Section 3 introduces the survey taxonomy and formal framing. Section 4 reviews major optimization approaches for SE minimization. Section 5 discusses learning-based formulations, including graph-learning and reinforcement-learning variants. Section 6 examines evaluation practice through our benchmark and reproducibility lens. Section 7 and Section 8 widen the view to implications, limitations, and future directions, and Section 9 concludes.

2 Related Work

SE belongs to a broader family of graph-complexity, graph-clustering, and graph-compression methods, but it differs from most classical objectives in treating hierarchical organization as the primary object to be decoded. The closest reference points for understanding its role are modularity-based community detection, flow-based compression, spectral partitioning, and modern deep graph clustering. These neighboring lines of work provide useful comparisons because they often pursue similar structural goals while relying on different assumptions about what counts as a good partition.

Modularity-based community detection. Modularity compares observed within-community edge density against a configuration-model null graph [33]. Louvain and Leiden-style optimization make this objective highly scalable, and Leiden further improves partition connectivity and refinement stability [57]. These methods remain strong default baselines for topology-driven clustering, but they are fundamentally flat and are known to exhibit resolution effects. SE differs in that it evaluates the coding cost induced by an encoding tree, so the partition is interpreted as an information hierarchy rather than only a density surplus.

Flow-based compression. Flow-based compression methods such as the map equation and Infomap encode random-walk trajectories using module codebooks [47]. They are conceptually close to SE because both explain graph organization through compression, but the object being compressed is different. Infomap compresses flows generated by random walks, whereas SE compresses structural uncertainty under graph volumes, cuts, and tree refinement. This distinction becomes important when the graph is interpreted as an information system rather than an explicit navigation process.

Spectral and cut-based methods. Spectral and cut-based methods use eigenvectors of graph Laplacians to expose low-conductance partitions [49, 34]. They provide powerful relaxations for balanced graph partitioning and remain important for initialization, diagnostics, and theoretical comparison. In most forms, however, spectral objectives are flat and externally parameterized by the target number of clusters. SE can connect to spectral information, including the literature around von Neumann entropy, but its target remains a minimum-description hierarchy rather than an eigenspace partition alone.

Deep graph clustering and pooling. Deep graph clustering and pooling methods such as DiffPool, MinCutPool, and DMoN learn soft assignments inside graph neural network architectures [78, 1, 58]. They are effective when node attributes and task supervision provide extra semantic signal, but their objectives are often architectural regularizers rather than standalone measures of structural organization. Recent SE-based neural methods, including LSENet, SEP-style pooling, and DeSE, help bridge this gap by making structural criteria usable inside learned models [66, 53, 94]. One goal of the present survey is to place these neural methods on the same conceptual map as the original combinatorial SE algorithms and their domain-specific extensions.

Taken together, these neighboring paradigms clarify what is distinctive about SE. It is not simply another flat graph-clustering score, nor only a regularizer for graph neural networks. Its central claim is that graph organization should be understood through the coding efficiency of hierarchical structure, and that this perspective can be used both analytically and algorithmically across a wide range of graph intelligence tasks.

3 Formal Foundations and Taxonomy

3.1 Basic Mathematical Constructs

3.1.1 Encoding Tree of a Finite Set

Let A be a finite set representing the entities in a complex system (e.g., nodes in a network or data points in a dataset). A natural model for the structure of such a system is a *hierarchical abstraction* in which elements are recursively grouped into nested subsets.

Definition 3.1 (Encoding Tree). *Given a finite set A , an encoding tree of A is a rooted tree T such that: (1) the root λ corresponds to the entire set $T_\lambda = A$; (2) for every non-leaf node $\alpha \in T$ with children $\beta_0, \dots, \beta_l$ the corresponding subsets form a partition $T_\alpha = \bigsqcup_{j=0}^l T_{\beta_j}$; (3) each leaf $\gamma \in T$ corresponds to a singleton $T_\gamma = \{a\}$.*

The *dimension* of an encoding tree is its depth. The extreme case—an encoding tree of depth 1 where the root is directly partitioned into singleton leaves—yields the **1-dimensional SE**, which reduces to Shannon entropy.

3.1.2 Structural Entropy

Consider a non-negative, irreducible matrix $A \in \mathbb{R}_{\geq 0}^{n \times n}$ representing an information system over n elements indexed by $V = \{1, \dots, n\}$. Let $\pi = (\pi_1, \dots, \pi_n)$ be the stationary distribution of the Markov chain derived from A by row normalization, with transition probability $p_{xy} = \pi_x b_{xy}$ where $b_{xy} = a_{xy} / \sum_j a_{xj}$.

For a subset $X \subseteq V$ define:

$$V_X = \sum_{x \in X} \pi_x, \quad p_X = \sum_{y \notin X} \sum_{x \in X} p_{yx}, \quad q_X = \sum_{x \in X} \sum_{y \notin X} p_{xy}.$$

Definition 3.2 (Structural Entropy under a Tree). *The structural entropy of A under encoding tree T is*

$$\mathcal{H}^T(A) = - \sum_{\substack{\alpha \in T \\ \alpha^- \neq \lambda}} p_\alpha \log_2 \frac{V_\alpha}{V_{\alpha^-}},$$

where α^- denotes the parent of α in T .

Definition 3.3 (Structural Entropy of a System). *The structural entropy of an information system is*

$$\mathcal{H}(A) = \min_T \mathcal{H}^T(A).$$

The minimizing tree T^* encodes the system's most efficient information-theoretic abstraction—its natural hierarchical community structure [16].

3.2 Relationship to Other Graph-Complexity Measures

Three major graph-complexity measures are related to SE:

Modularity [33]. Modularity Q measures the fraction of edges within communities minus the expected fraction under a null model. While modularity seeks to maximize within-community density relative to a random graph, SE minimization seeks to *minimize informational uncertainty* within a hierarchical partition. Modularity is resolution-limited (it cannot detect communities smaller than a scale-dependent threshold); SE is not subject to the same limitation.

Map Equation [47]. Infomap minimizes the expected description length of a random walk under a two-level codebook. SE generalization replaces the two-level restriction with an arbitrary tree, and replaces flow coding with a volume-and-cut measure. Formally, [31] proved that for any undirected graph $\mathcal{H}(A) \leq L^*(A) \leq \mathcal{H}(A) + \gamma$, where L^* is the optimal map-equation description length and γ depends only on the graph's degree sequence.

Von Neumann Entropy (VNE). VNE is the entropy of the normalized graph Laplacian's eigenvalue distribution. [31] established that $\mathcal{H}(A) - \text{VNE}(A) \in [0, \gamma]$, providing the first formal spectral interpretation of SE.

3.3 Taxonomy of SE Research

Figure 1 provides a high-level taxonomy organizing SE research into four pillars—*Theory*, *Computational Algorithms*, *Learning Methods*, and *Cross-Domain Applications*—each containing major sub-themes and representative works. We follow this taxonomy throughout the survey.

4 Theoretical Foundations and Computational Algorithms

4.1 Foundational Theoretical Works

Foundation Theory. The concept of SE was introduced by [16] as a principled, information-theoretic framework for quantifying the complexity of a graph’s structure. The seminal work defined SE as the minimum code length required to describe a network’s hierarchy under a random walk process, bridging graph theory with Shannon’s information principles. It established that minimizing SE reveals a graph’s intrinsic multi-scale organization. Building on this, [14] significantly broadened the theoretical landscape by situating SE within the *mathematical principles of the information world*, formulating fundamental information laws and the Information Science Principles of AI (ISPAI), providing axiomatic foundations for a comprehensive structural information theory.

Theoretical Analysis. Subsequent research has elucidated SE’s properties and forged connections with established graph-theoretic concepts. The “togetherness” metric [96] uses SE reduction to quantify the efficacy of network integration when adding inter-community edges. [31] established a significant theoretical link between SE and spectral graph theory: for any undirected graph, the difference between SE and the von Neumann graph entropy (VNE) is bounded within the interval $[0, \gamma]$, formally connecting SE to spectral properties of the graph Laplacian. [71] proposed incremental computation paradigms for dynamic graphs, enabling efficient encoding tree updates as the network evolves.

Extensions. The original formulation targets simple graphs; significant research has extended SE to more complex network structures. For directed graphs, [76] introduced a localized 2D SE via a flow-based algorithm, while [92] proposed a random walk-based entropy (RWE) with a proven theoretical connection to VNE. Multi-relational SE [2] uses a random surfing model that jointly optimizes node and relation selection. The *infected tree entropy* (ITE) [93] specializes SE for source localization in information cascades. SE-guided probabilistic coding [12] regularizes embeddings in natural language understanding tasks.

4.2 Computational Algorithms

Table 1 summarizes all known SE minimization algorithms. We organize them by optimization strategy.

Heuristic Methods. deDoc [20] pioneered SE minimization for genomic structure detection with a greedy merging strategy. SEAT [4] adapts this for single-cell omics by




Figure 1: Taxonomy of Structural Entropy research for the TGINA journal version (2016–2026), following the vertical visual organization of the IJCAI conference taxonomy [51] while adding journal-version categories and citations.

Table 1: Summary of SE minimization algorithms. n : nodes; m : edges; t : iterations; k : flat clusters; h : arbitrary tree depth.

Algorithm	Levels	Constraint	Time Complexity	Code	Lang
HCSE [38]	h	Unconstrained	$O(n^2)$	Link	Python
deDoc [20]	2/3	Unconstrained	$O(n \log^2 n)$ or $O(n^2)$	Link	Java
deDoc2 [21]	h	Unconstrained	$O(n^2)$	Link	Java
SSE [80]	h	Label, Pairwise	$O(h(m \log n + n))$	Link	Python
SSSE [82]	h	Label, Pairwise	$O(n \log^2 n + nkt)$	Link	Python
SuperTAD [97]	h	Contiguity	$O(n^4 k^2 h)$	Link	C++
SuperTAD2 [27]	h	Contiguity	$O(n^4 k^2 h)$	Link	C++
PYSEAT [4]	h	Contiguity	$O(n \log n)$	Link	Python
LSENet [53]	2	Unconstrained	$O(nm)$	Link	Python
DeSE [94]	2	Unconstrained	$O(nm)$	—	Python
CoDeSEG [67]	2	Overlapping	$O(nt)$	Link	C++

applying SE minimization on cell–cell graphs to reveal nested subpopulations. HCSE [38] is a widely used, general-purpose hierarchical clustering framework that recursively isolates the sparsest levels of a hierarchy through SE-guided optimization, automatically determining hierarchy depth without hyperparameters.

Constrained Optimization. SuperTAD [97] combines interval dynamic programming with SE minimization under *spatial contiguity* constraints. SuperTAD2 [27] introduces matrix discretization and approximation strategies achieving $\approx 10\times$ speedup. deDoc2 [21] extends the contiguity approach to single-cell analysis to resolve chromatin hierarchies across the cell cycle. SSE [80] unifies *semi-supervised* constraints (must-link / cannot-link pairwise constraints and label constraints) by reformulating SE to penalize violations during encoding tree construction. SSSE [82] scales this to massive datasets via graph sampling and incremental cluster insertion, reducing time complexity from quadratic to nearly linear.

Differentiable / Gradient-Based Methods. LSENet [53] creates a differentiable SE formulation in hyperbolic space, enabling cluster-free hierarchical partitioning via backpropagation. DeSE [94] pioneers a *soft assignment structural entropy* allowing co-design of graph topology and cluster assignments in an end-to-end model. CoDeSEG [67] offers a game-theoretic formulation for near-linear-time overlapping community detection.

Complexity and Trade-off Analysis. Heuristic approaches offer scalability and generality but are limited in their ability to incorporate domain knowledge. Constrained methods integrate domain-specific knowledge at the cost of broader applicability. Differentiable methods bridge SE and deep learning but introduce hyperparameter sensitivity and potential interpretability concerns. A key open challenge is hybridizing these paradigms to achieve an optimal balance of scalability, accuracy, and interpretability, especially for dynamic or heterogeneous networks (see Section 8).

5 Structural Entropy in Graph Learning and Reinforcement Learning

5.1 Graph Learning Paradigms

SE has emerged as a powerful information-theoretic tool for graph representation learning, offering a principled, quantitative measure of a graph’s inherent hierarchical organization.

5.1.1 Graph Pooling

In graph-level tasks, pooling operations are essential for downsampling nodes to create compact representations of entire graphs. Traditional methods often suffer from information loss by applying stepwise local pooling, which destroys the larger community structure. *SE Pooling* (SEP) [66] addresses this by using global SE minimization to generate all cluster assignments simultaneously, preserving the multi-level community map. *Hi-PART* [45] constructs Hierarchical Partition Trees using SE, explicitly modelling the nested, multi-scale nature of graph structures and surpassing the expressive power of the standard Weisfeiler-Lehman test. *EDEN* [25] further extends SE-based pooling to dynamic and evolving graphs.

5.1.2 Structure Augmentation

Graph structure augmentation underpins modern self-supervised learning on graphs. Random augmentations risk discarding crucial information; SE provides a principled way to guide this process by identifying and preserving the graph’s intrinsic low-entropy structure. *SEGA* [65] constructs a stable anchor view by finding an encoding tree with minimal SE, filtering noise while maintaining the essential community hierarchy for contrastive learning. *USER* [64] incorporates SE directly into its objective function to learn an “innocuous graph” robust to perturbations. *SE-GSL* [102] uses a three-stage optimization: maximizing 1D SE to enhance information content, minimizing high-dimensional SE to abstract the graph into hierarchical trees, then using deduction SE to guide adaptive graph reconstruction. *SeGSL* [7] extends this by using SE-derived minimal-tree structures to guide multi-view fusion based on community influence.

5.1.3 Graph Kernels and Embedding Dimensions

HAGK [70] leverages SE to generate multi-level, abstract representations of each graph’s intrinsic organization for use in a graph kernel that measures similarity between entire graphs. *MGEDE* [73] calculates a high-level SE from multi-order adjacency matrices and minimizes it alongside attribute entropy, explicitly determining the most suitable number of embedding dimensions.

5.2 Reinforcement Learning Integration

Integrating SE into reinforcement learning (RL) equips agents with a mechanism to perceive and decompose complex problems, moving beyond flat state-action spaces toward a structured understanding of their environment.

5.2.1 State and Action Abstraction

The *SISA* [88] framework uncovers hierarchical state structure by optimizing an encoding tree from a state graph using SE, creating an efficient and task-relevant state representation. *SIRD* [83] reframes role discovery in multi-agent RL as hierarchical action space clustering via SE minimization on an action graph.

5.2.2 Intrinsic Motivation and Exploration

SI2E [85] formulates an intrinsic reward by maximizing value-conditional SE: agents are encouraged to explore states where their actions lead to the greatest reduction in structural uncertainty. *SIDM* [91] utilizes SE at multiple stages to structure the entire decision-making and skill-discovery pipeline. *COLLAB-MARL* [52] uses SE to quantify interaction complexity in multi-agent systems, revealing emergent coalition structures by minimizing the SE of a dynamic agent interaction graph.

5.2.3 Discussion

SE integration into RL provides mathematical tools to build a *cognitive architecture* for agents: parsing their world into a hierarchy of concepts, skills, and goals. *SISA* and *SIRD* demonstrate how SE can automate the discovery of abstractions that have historically required significant human engineering. *SI2E* shows that exploration need not be random—it can be a deliberate process of *information seeking* aimed at resolving structural uncertainty. Key future directions include scaling these methods to partially observable environments, meta-RL, and curriculum learning.

5.3 Cross-Domain Applications

5.3.1 Bioinformatics

SE-based genomic analysis began with cancer subtyping [19], correlating entropy-minimized gene expression patterns with clinical outcomes. The first SE-based TAD detector, *deDoc* [20], identified large-scale topologically associating domains from Hi-C data. *SuperTAD* [97] employed dynamic programming to resolve hierarchical chromatin domains; *SuperTAD2* [27] significantly accelerated this via matrix discretization. At single-cell resolution, *SEAT* [4] minimizes global uncertainty in cell–cell graphs to improve pseudo-time inference. *deTOKI* [24] analyzes regulatory dynamics at the single-cell level, and *deDoc2* [21] resolves cell-cycle-dependent chromatin hierarchies.

5.3.2 Transportation and Geoscience

MultiSPANS [103] improves traffic forecasting by using SE minimization to optimize spatial attention mechanisms, generating hierarchical encoding trees of road networks. The “Scientific X-ray” method [23] constructs “idea trees” from citation networks to identify high-potential research directions. *KQI* [63] redefines knowledge quantification by measuring hierarchical disorder differences in citation networks, producing a Knowledge Quantification Index that identifies overlooked but influential works and Nobel Prize-winning topics.

5.3.3 Social Networks

UnDBot [41] offers unsupervised bot detection via heterogeneous SE on multi-relational graphs. *SEBot* [72] integrates SE with multi-view contrastive learning for adversarially robust detection. *SIASM* [84] models how social bots optimally select followers using conditional SE minimization. *SI2AF* [86] evaluates news-detection system robustness using an SE-based influence metric that combines content relevance with community structure. For social event detection, *HISEvent* [3] incrementally builds message graphs with 1D SE and detects events using 2D SE optimization; *ADP-SEMEvent* [75] integrates adaptive differential privacy; *HyperSED* [79] uses hyperbolic-space differentiable SE-guided partitioning.

5.3.4 Pattern Recognition

In *visual computing*, *SLED* [81] detects skin lesions by minimizing the entropy of multi-scale superpixel graphs. *SIT-HSS* [68] performs hierarchical superpixel segmentation via iterative SE-guided merging.

In *natural language processing*, *HiTIN* [101] injects the syntactic hierarchy of sentences via SE into graph-based representations for hierarchical text classification. *MASGCN* [13] applies an SE-based loss to guide a GCN for aspect-based sentiment analysis.

In *speech processing*, *SECodec* [61] applies SE to cluster feature nodes for efficient compressed speech representation. *SEVC* [62] uses 2D SE to group reference speech frames into semantic clusters for high-fidelity voice conversion.

6 Reproducibility Benchmark

6.1 Experimental Protocol

To complement our survey with empirical evidence, we present a reproducibility benchmark using the `glass-jax` companion library—a JAX-based differentiable graph clustering framework that implements soft versions of Modularity, the Map Equation, and 2D Structural Entropy (see Appendix A for library details). All results are reported as **mean $\pm$ 95% confidence interval** over 5 random seeds.

Datasets.

- **Karate**: Zachary’s Karate Club ($N=34$, $K=2$).
- **Caveman(10 $\times$ 10)**: Connected Caveman graph of 10 cliques of size 10 ($N=100$, $K=10$).
- **SBM (Clean)**: Stochastic Block Model ($N=150$, $K=3$, $P_{\text{in}}=0.30$, $P_{\text{out}}=0.05$).
- **SBM (Noisy)**: Stochastic Block Model ($N=150$, $K=3$, $P_{\text{in}}=0.15$, $P_{\text{out}}=0.08$; signal near the community-detection threshold).

Algorithms compared. We include two classes of baselines:

Non-SE baselines (community detection):

1. **Louvain** [33]: official `python-louvain` implementation.
2. **Leiden**: improved Louvain with guaranteed connectivity.
3. **Infomap** [47]: official C++ Python wrapper.
4. **Spectral**: normalized spectral clustering with k -means on the top- K Laplacian eigenvectors.

SE-based methods:

5. **deDoc** [20]: the founding SE algorithm; greedy agglomerative merging that minimises 2D structural entropy at each step (`agglomerative_se_clustering` in the companion library). The reference implementation is $O(N^3)$; results are reported only for Karate ($N=34$); larger datasets are marked $-$.
6. **HCSE** [38]: hierarchical SE clustering. Results here use the companion library’s `hierarchical_se_clustering`, which builds a coding tree from multistart base labels and extracts a flat cut.
7. **SEClust-Auto***: `multistart_se_heuristic` (8 restarts, free K) from the companion library.
8. **SEClust-CK***: `constrained_k_multistart`; multistart SE local search with K fixed to the ground-truth value.
9. **SEClust-ML***: `multilevel_se_clustering`; multi-level coarsen-and-refine SE optimizer (free K).
10. **SEP** [66]: greedy 2D SE pooling algorithm.
11. **LSENet** [53]: differentiable SE via neural network projection.
12. **Glass-Mod**: soft Modularity maximization (`glass-jax`).
13. **Glass-Map**: soft Map Equation minimization (`glass-jax`).
14. **Glass-SE**: soft 2D SE minimization (`glass-jax`).

* From the `glass-jax` companion library (this survey).

All `glass-jax` methods use a multi-start `vmap` engine: four parallel trajectories (1 spectral + 3 random initializations) over 500 Adam steps ($\eta = 0.05$) with temperature annealing from $\tau = 1.0$ to $\tau = 0.01$.

Algorithms excluded from this benchmark. Three SE methods from Table 1 require a fundamentally different evaluation protocol and are therefore excluded. (i) **SSE** [80] is *semi-supervised*: it takes a subset of ground-truth node labels as input, so evaluating it in an unsupervised protocol would be misleading. (ii) **CoDeSEG** [67] detects *overlapping* communities; the standard ARI/NMI metrics for disjoint partitions are not applicable, and the method requires overlapping-specific metrics (F1-overlap, NMI-overlap). (iii) **DeSE** [94] jointly trains topology and assignments end-to-end; no standalone graph-partition interface or public code is available at the survey submission date. **Super-TAD** [97] additionally requires a Hi-C contiguity constraint not present in generic graph benchmarks. **deDoc** is included above for Karate only; the reference implementation is $O(N^3)$ and infeasible for $N > 50$.

6.2 Results: Clustering Accuracy

Table 2 shows ARI and NMI for all algorithms on all datasets.




Figure 2: Representative ARI comparison across synthetic and small real-world topological benchmarks. Glass-SE is competitive with strong baselines on clean and noisy topology, while feature-augmented LSENet is strongest when semantic features align with labels.

Table 2: Clustering accuracy (ARI and NMI, mean $\pm$ 95% CI over 5 seeds). Higher is better. Bold: best per dataset. †: SE-based method. *: companion library (this survey). -: method infeasible at this scale ($O(N^3)$, $N>50$).

Dataset	Algorithm	ARI	NMI
Karate ($N=34$)	Louvain	0.500 $\pm$ 0.010	0.597 $\pm$ 0.003
	Leiden	0.500 $\pm$ 0.010	0.597 $\pm$ 0.003
	Infomap	0.684 $\pm$ 0.000	0.691 $\pm$ 0.000
	Spectral	0.620 $\pm$ 0.040	0.650 $\pm$ 0.030
	deDoc†	0.434 $\pm$ 0.000	0.571 $\pm$ 0.000
	HCSE†	0.882$\pm$0.000	0.837$\pm$0.000
	SEClust-Auto†*	0.337 $\pm$ 0.052	0.516 $\pm$ 0.029
	SEClust-CK†*	0.882$\pm$0.000	0.837$\pm$0.000
	SEClust-ML†*	0.393 $\pm$ 0.069	0.546 $\pm$ 0.041
	SEP†	0.882$\pm$0.000	0.837$\pm$0.000
	LSENet†	0.882$\pm$0.000	0.837$\pm$0.000
	Glass-Mod	0.882$\pm$0.000	0.837$\pm$0.000
Caveman ($N=100$)	Glass-Map	0.444 $\pm$ 0.278	0.458 $\pm$ 0.273
	Glass-SE†	0.882$\pm$0.000	0.837$\pm$0.000
	Louvain	1.000$\pm$0.000	1.000$\pm$0.000
	Leiden	1.000$\pm$0.000	1.000$\pm$0.000
	Infomap	1.000$\pm$0.000	1.000$\pm$0.000
	Spectral	1.000$\pm$0.000	1.000$\pm$0.000
	deDoc†	—	—
	HCSE†	0.421 $\pm$ 0.000	0.733 $\pm$ 0.000
	SEClust-Auto†*	1.000$\pm$0.000	1.000$\pm$0.000
	SEClust-CK†*	0.889 $\pm$ 0.000	0.969 $\pm$ 0.000
	SEClust-ML†*	1.000$\pm$0.000	1.000$\pm$0.000
	SEP†	1.000$\pm$0.000	1.000$\pm$0.000
SBM Clean ($N=150$)	LSENet†	0.956 $\pm$ 0.048	0.988 $\pm$ 0.013
	Glass-Mod	0.852 $\pm$ 0.040	0.956 $\pm$ 0.014
	Glass-Map	0.786 $\pm$ 0.000	0.929 $\pm$ 0.000
	Glass-SE†	0.893 $\pm$ 0.057	0.969 $\pm$ 0.018
	Louvain	1.000$\pm$0.000	1.000$\pm$0.000
	Leiden	1.000$\pm$0.000	1.000$\pm$0.000
	Infomap	0.996 $\pm$ 0.007	0.994 $\pm$ 0.010
	Spectral	1.000$\pm$0.000	1.000$\pm$0.000
SBM Clean ($N=150$)	deDoc†	—	—
	HCSE†	0.984 $\pm$ 0.013	0.978 $\pm$ 0.017
	SEClust-Auto†*	0.562 $\pm$ 0.039	0.737 $\pm$ 0.020

Dataset	Algorithm	ARI	NMI
SBM Noisy ($N=150$)	SEClust-CK ^{†*}	0.992±0.014	0.988±0.021
	SEClust-ML ^{†*}	0.611±0.052	0.767±0.022
	SEP [†]	0.949±0.036	0.932±0.048
	LSENet [†]	1.000±0.000	1.000±0.000
	Glass-Mod	0.996±0.007	0.994±0.010
	Glass-Map	0.682±0.232	0.764±0.172
	Glass-SE [†]	0.996±0.007	0.994±0.010
	Louvain	0.080±0.031	0.129±0.022
	Leiden	0.083±0.028	0.131±0.020
	Infomap	0.000±0.000	0.000±0.000
	Spectral	0.095±0.045	0.140±0.038
	deDoc [†]	—	—
	HCSE [†]	0.055±0.036	0.058±0.033
	SEClust-Auto ^{†*}	0.051±0.017	0.119±0.029
	SEClust-CK ^{†*}	0.085±0.028	0.088±0.023
	SEClust-ML ^{†*}	0.050±0.035	0.120±0.046
	SEP [†]	0.064±0.027	0.078±0.026
	LSENet [†]	0.118±0.056	0.116±0.048
	Glass-Mod	0.136±0.096	0.133±0.081
	Glass-Map	0.090±0.062	0.138±0.054
	Glass-SE[†]	0.138±0.106	0.135±0.091

Note on deDoc: The reference agglomerative implementation is $O(N^3)$ per run; for Karate ($N=34$) each seed takes ~ 5.6 s. For $N>50$ the runtime is prohibitive and results are marked —. An $O(m \log N)$ online update (as in the original paper) would resolve this limitation. **Note on HCSE:** The companion library’s `hierarchical_se_clustering` builds a binary coding tree (`target_k=2`); for $K>2$ grounds truth the tree undercounts communities, reducing ARI (e.g., Caveman $K=10$: ARI 0.421). The original Pan et al. [38] construction is depth-unlimited and achieves higher accuracy. **Note on SEClust-Auto vs SEClust-CK:** SEClust-Auto finds K automatically; SEClust-CK uses ground-truth K . Because the SBM signal-to-noise is low, even with the correct K all heuristics score at the detection threshold on SBM-Noisy.

6.3 Results: Objective Evaluators

Table 3 cross-evaluates every algorithm’s output against all three clustering objectives.

6.4 Results: Real-World Attributed Graphs

Table 4 reports results on the two standard attributed citation networks used in the LSENet paper [53]. Crucially, the ground-truth partition in these datasets is *semantic* (paper topics) rather than topological (citation density).

Topology-only vs. feature-augmented. On feature-rich datasets, feature-augmented LSENet outperforms topology-only Glass-SE in ACC (Cora: $0.387 > 0.274$; Citeseer: $0.403 > 0.252$), because the ground-truth classes (academic subjects) are poorly aligned with citation-graph topology. Louvain dominates on NMI, reflecting that the citation structure still carries significant mutual information with ground truth, even if it does not fully separate classes. This is consistent with prior observations [53, 8] that topology-only SE objectives are best suited to graphs where cluster boundaries coincide with connectivity bottlenecks. Incorporating node features via a GNN encoder — as in LSENet — is the natural path to extending structural entropy minimization to attributed graphs.

6.5 Analysis and Takeaways

Noisy graph superiority. On the challenging **SBM (Noisy)** dataset (signal near the community-detection threshold), Glass-SE achieved ARI 0.138, the highest among all tested methods. All other methods — Louvain (0.080), Leiden (0.083), Infomap (0.000), Spectral (0.095), HCSE (0.055), SEClust-CK (0.085), SEClust-ML (0.050), SEP (0.064), LSENet (0.118), Glass-Mod (0.136) — fall at or below Glass-SE on this hardest regime. This demonstrates the resilience of direct soft-logit optimization via `vmap` under high noise; the continuous relaxation effectively averages over multiple near-optimal partitions.

Clean graphs: K -constrained SE dominates. On **SBM (Clean)** and **Karate**, SEClust-CK (ARI 0.992 and 0.882) and HCSE (0.984 and 0.882) match or nearly match the best achievable performance, confirming that SE minimization with correct K is a reliable and accurate clustering criterion when the signal is strong.

Caveman: free- K multi-level wins. On the Caveman graph (10 perfectly-disjoint cliques), SEClust-ML and all discrete baselines achieve ARI 1.000, while HCSE drops to 0.421. The failure of the binary coding tree (HCSE) on $K=10$ reflects the mismatch between binary tree structure and a K -way ground truth. Free- K multi-level search naturally discovers all 10 components.

Caveman limitation for continuous methods. The only setting where discrete methods (ARI 1.000) clearly exceeded Glass-SE (0.893) was the Caveman graph, whose perfectly disjoint cliques present a degenerate topology. Because there are literally zero inter-clique edges, no gradient signal propagates to push cluster assignments apart; continuous relaxation struggles to represent absolute zero probabilities. An L1 sparsity penalty or feature projection layer (similar to LSENet’s MLP) resolves this, but is unnecessary for real-world, dense, noisy graphs.

K -constraint artifact. On Karate, Louvain achieves a lower (better) raw SE objective (3.323) than Glass-SE (3.731) because Louvain freely optimizes K (choosing $K=4$) while accuracy-benchmarked Glass-SE is constrained to $K=2$. When Glass-SE is allowed to optimize freely over $K=8$, it finds SE 3.245, beating Louvain’s global minimum.

Companion library. All experimental code and configurations are released in `glass-jax` at <https://github.com/SELGroup/glass-jax>. See Appendix A for reproduction instructions.

Runtime and scalability. The benchmark also highlights a practical scalability divide. Discrete baselines such as Louvain and Leiden remain the fastest choices for large sparse graphs because they operate directly on sparse neighborhoods. Reference deDoc-style agglomeration is cubic in its naive form and is therefore reported only for Karate; in our implementation each Karate seed takes about 5.6 seconds, and $N>50$ runs are marked infeasible. The differentiable Glass objectives are optimized by dense soft assignments over 500 Adam steps and four vectorized starts, making them most appropriate for medium-scale graphs or neural latent graphs where gradient flow matters. Scaling differentiable SE to million-node graphs will require sparse kernels, mini-batch approximations, or coarsen-and-refine objectives, which we identify as a major open problem in Section 8.

Benchmark scope and open problem O7. This benchmark covers representative topology-only methods on synthetic graphs and real-world attributed graphs, but several important method classes remain outside its scope. Hierarchical SE methods (HCSE hierarchy depth, Paris, Agglomerative) require a dedicated dendrogram quality metric (e.g., Dasgupta cost), which we defer to future work (see Open Problem O7, §8). Scalability evaluation on million-node graphs (LFR benchmarks, ogbn-arxiv) is likewise deferred, pending `glass-jax` sparse-matrix support. The authors encourage the community to contribute standardized SE benchmark results to <https://github.com/SELGroup/glass-jax/benchmarks>.

7 Discussion

The decade of work surveyed above suggests that SE is best understood not as a single clustering score, but as a unifying principle for graph intelligence: systems become intelligible when their uncertainty can be reduced by a compact hierarchy. This interpretation explains why SE appears in apparently distant settings, including graph pooling, chromatin-domain discovery, event detection, social bot analysis, and reinforcement learning. In each case, the method searches for a structural abstraction that preserves the dominant organization while suppressing incidental variation.

For TGINA’s focus on graph intelligence and network applications, SE is especially useful because it connects three concerns that are often handled separately. First, it is a *measure*: the entropy value summarizes how ordered or disordered a network is under a candidate hierarchy. Second, it is an *optimizer*: minimizing the measure constructs the hierarchy itself. Third, it is a *learning signal*: differentiable variants allow the same principle to shape GNN encoders, pooling layers, and decision abstractions. This measure–optimizer–learning-signal triad is the core reason SE deserves a full journal treatment rather than a short method catalogue.

At the same time, the benchmark results caution against overclaiming. Topology-only SE objectives are most appropriate when communities correspond to structural bottlenecks. On attributed citation networks, semantic labels may be better explained by text features than by citation topology, and feature-augmented models such as LSENet can therefore dominate in accuracy. Conversely, on noisy synthetic topologies, differentiable SE remains competitive with or stronger than standard baselines. The practical lesson is that SE should be matched to the source of structure: pure topology for latent abstraction and network organization, feature-aware SE for semantic graph learning, and constrained SE for domain-specific structures such as genomic intervals.

Finally, SE offers a promising language for trustworthy graph intelligence in LLM and autonomous-agent systems. GraphRAG pipelines, tool-use traces, memory graphs, and multi-agent interaction networks all produce evolving relational structures whose reliability depends on whether useful organization can be separated from noisy connectivity. SE can quantify that organization, detect disorder, and potentially guide the construction of stable retrieval or decision hierarchies. This connection remains early, but it gives a concrete path from classical structural information theory to the emerging graph intelligence problems that motivate the TGINA venue.

8 Open Problems and Future Directions

Despite a decade of rapid progress, several fundamental challenges remain open:

O1. Unified Framework for Complex Network Types

Existing SE variants have been developed independently for directed graphs, hypergraphs, multi-relational graphs, and dynamic networks. A pressing open problem is the development of a *unified SE framework* that handles these structures under a single formalism. This would require extending the encoding-tree formulation to support non-standard volume and cut measures, and establishing formal connections between the specialized variants.

O2. Scalable SE Minimization

The dominant heuristic SE algorithms scale as $O(n^2)$, limiting applicability to graphs with hundreds of thousands of nodes. Although SSSE achieves nearly linear complexity under a sampling approximation, theoretical guarantees on the approximation quality remain weak. Developing provably good approximation algorithms for SE minimization, analogous to those for modularity and the map equation, is an important open direction.

O3. Differentiable SE for Deep Learning Integration

Current differentiable SE formulations (LSENet, DeSE, Glass-SE) are restricted to 2-level (flat) trees. Extending differentiable SE to arbitrary-depth hierarchical trees—while maintaining well-defined gradients and computational tractability—is an open problem. This would enable end-to-end learning of hierarchical representations in GNNs and Transformers.

O4. SE and Large Language Models

The recent success of LLMs raises the question of whether SE can be used to analyse or improve the latent graph structures induced by LLM attention patterns and knowledge graphs. Preliminary connections have been drawn via SE-based knowledge quantification [63] and graph-augmented reasoning [26], but a principled SE-based theory of LLM structure is absent.

O5. Theoretical Connections to Other Complexity Measures

While [31] established a bound between SE and VNE, the exact relationship between SE and other information-theoretic graph measures (e.g., Kolmogorov complexity, Fisher information, persistent homology) is not well understood. A deeper theoretical investigation would both strengthen SE’s foundations and enable new algorithmic insights.

O6. SE in Partially Observable Environments

SE-based RL methods (SISA, SIRD, SI2E) have been developed for fully observable settings. Extending SE to partial observability—where the agent’s state graph is uncertain or must be inferred—is an important frontier for real-world applicability.

O7. Benchmark Standardization

A major bottleneck for the field is the absence of a standardized evaluation protocol. Different papers use different datasets, train–test splits, and evaluation metrics, making direct comparison difficult. We call for the community to adopt a standard benchmark suite (as partially initiated in Section 6) with open, reproducible implementations.

O8. SE for Fairness and Privacy

The use of SE to deliberately obscure community structure [32] suggests promising applications in *privacy-preserving graph publishing* and *fairness-aware graph learning*. A principled theory connecting SE-based obfuscation to differential privacy guarantees remains an open problem.

9 Conclusion

This survey has systematically charted the landscape of Structural Entropy research over the decade since its introduction in 2016. We have consolidated its theoretical underpinnings, computational methodologies, and burgeoning applications across AI, presenting the field’s first comprehensive synthesis.

Our review reveals that SE offers a distinct and powerful lens for artificial intelligence: it shifts the focus from local, pairwise interactions to the global, hierarchical organization of complex systems. We have detailed how minimizing SE serves as a constructive principle to uncover a system’s most efficient information-theoretic abstraction—a capability that has been effectively harnessed in graph learning for pooling and structure augmentation, in reinforcement learning for hierarchical state–action abstraction and exploration, and in diverse application domains from genomics to speech processing.

The reproducibility benchmark in Section 6 demonstrates that differentiable SE objectives, as implemented in the `glass-jax` library, match or exceed state-of-the-art methods on realistic, noisy graph topologies while offering the key advantage of end-to-end gradient-based optimization.

In conclusion, Structural Entropy provides more than a new set of tools: it offers a principled way of thinking about information and structure in a hierarchical world. By providing a rigorous method to quantify and optimize organizational uncertainty within data, SE holds the potential to guide the development of AI systems that are not only more powerful but also more interpretable, robust, and theoretically grounded. The eight open problems identified in Section 8 chart a research agenda for the next decade.

Differences from the IJCAI-2025 Conference Version. This journal article substantially extends [51] in the following ways:

1. A new, self-contained *formal foundations* section with precise definitions and comparative analysis against modularity, map equation, and von Neumann entropy.
2. An extended *algorithm catalogue* (Table 1) covering all known SE minimization methods with unified complexity notation.

3. A new *reproducibility benchmark* (Section 6) with empirical comparisons over multiple datasets and random seeds.
4. Expanded *application coverage*: speech processing (SECodec, SEVC), geoscience knowledge quantification (KQI), and updated 2025–2026 literature.
5. A structured *open problems* section (Section 8) absent from the conference version.
6. Release of `glass-jax` as a companion reproducibility library.

Acknowledgments

This work has been supported in part by NSFC through grants 62322202, 62441612 and 62432006, Local Science and Technology Development Fund of Hebei Province Guided by the Central Government of China through grant 246Z0102G, the “Pioneer” and “Leading Goose” R&D Program of Zhejiang through grant 2025C02044, National Key Laboratory under grant 241-HF-D07-01, Hebei Natural Science Foundation through grant F2024210008, National Key R&D Program of China through grant 2021YFB3500700, and State Key Laboratory of Complex & Critical Software Environment (CCSE-2024ZX-20).

References

- [1] Filippo Maria Bianchi, Daniele Grattarola, and Cesare Alippi. Spectral clustering with graph neural networks for graph pooling. In *Proceedings of the International Conference on Machine Learning*, 2020.
- [2] Yuwei Cao, Hao Peng, Angsheng Li, Chenyu You, Zhifeng Hao, and Yu Philip. Multi-relational structural entropy. In *Proceedings of the UAI*, 2024.
- [3] Yuwei Cao, Hao Peng, Zhengtao Yu, and S Yu Philip. Hierarchical and incremental structural entropy minimization for unsupervised social event detection. In *Proceedings of the AAAI*, 2024.
- [4] Lingxi Chen and Shuai Cheng Li. Incorporating cell hierarchy to decipher the functional diversity of single cells. *Nucleic Acids Research*, 2022.
- [5] Xiang Chen, Kun Yue, Daliang Liu, Wenjie Liu, Liang Duan, and Angsheng Li. Structural entropy guided meta-learning for few-shot node classification. *IEEE Transactions on Knowledge and Data Engineering*, 2026.
- [6] Sanjoy Dasgupta. A cost function for similarity-based hierarchical clustering. In *Proceedings of the STOC*, 2016.
- [7] Liang Duan, Xiang Chen, and et al. Structural entropy based graph structure learning for node classification. In *Proceedings of the AAAI*, 2024.
- [8] Santo Fortunato. Community detection in graphs. *Physics Reports*, 486(3–5):75–174, 2010.
- [9] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incen-

- tivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- [10] Yuxin Hou, Aoqi Zhang, Xianghua Zeng, Hao Peng, Angsheng Li, and Philip S. Yu. Structural entropy guided unsupervised graph out-of-distribution detection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 17258–17266, 2025.
 - [11] Junyu Huang, Qilong Feng, Jiahui Wang, Ziyun Huang, Jinhui Xu, and Jianxin Wang. Sec: More accurate clustering algorithm via structural entropy. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 12583–12590, 2024.
 - [12] Xiang Huang, Hao Peng, Li Sun, and et al. Structural entropy guided probabilistic coding. In *Proceedings of the AAAI*, 2024.
 - [13] Xiang Huang, Hao Peng, Shuo Sun, and et al. Multi-view attention syntactic enhanced graph convolutional network for aspect-based sentiment analysis. In *Proceedings of the DASFAA*, 2025.
 - [14] Angsheng Li. *Science of Artificial Intelligence: Mathematical Principles of Intelligence* (In Chinese). Science Press, Beijing, 2024.
 - [15] AngSheng Li, XueChen Li, YiCheng Pan, and Wei Zhang. Strategies for network security. *Science China Information Sciences*, 2015.
 - [16] Angsheng Li and Yicheng Pan. Structural information and dynamical complexity of networks. *IEEE Transactions on Information Theory*, 2016.
 - [17] Angsheng Li and Yicheng Pan. A theory of network security: Principles of natural selection and combinatorics. *Internet Mathematics*, 2016.
 - [18] Angsheng Li and Yicheng Pan. Structure entropy and resistor graphs. *arXiv preprint arXiv:1801.03404*, 2018.
 - [19] Angsheng Li, Xianchen Yin, and Yicheng Pan. Three-dimensional gene map of cancer cell types: Structural entropy minimisation principle for defining tumour subtypes. *Scientific reports*, 2016.
 - [20] Angsheng Li, Xianchen Yin, Bingxiang Xu, and et al. Decoding topologically associating domains with ultra-low resolution hi-c data by graph structural entropy. *Nature communications*, 2018.
 - [21] Angsheng Li, Guangjie Zeng, Haoyu Wang, Xiao Li, and Zhihua Zhang. Dedoc2 identifies and characterizes the hierarchy and dynamics of chromatin tad-like domains in the single cells. *Advanced Science*, 2023.
 - [22] Pu Li, Huafeng Li, Yafei Zhang, Yu Liu, and Wen Wang. Breaking degradation coupling: A structural entropy guided decoupled framework and benchmark for infrared enhancement, 2026.
 - [23] Qi Li, Xinbing Wang, and Luoyi and et al. Fu. Scientific x-ray: Scanning and quantifying the idea evolution of scientific publications. *Plos one*, 2022.

- [24] Xiao Li, Guangjie Zeng, Angsheng Li, and Zhihua Zhang. Detoki identifies and characterizes the dynamics of chromatin tad-like domains in a single cell. *Genome biology*, 2021.
- [25] Xunkai Li, Zhengyu Wu, Kaichi Yu, Hongchao Qin, Guang Zeng, Rong-Hua Li, and Guoren Wang. Toward data-centric directed graph learning: An entropy-driven approach. In *ICML*, 2025.
- [26] Yuhan Li, Zhixun Li, Peisong Wang, Jia Li, Xiangguo Sun, Hong Cheng, and Jeffrey Xu Yu. A survey of graph meets large language model: Progress and future directions. In *Proceedings of the IJCAI*, 2024.
- [27] Zhao Ling, Yu Wei Zhang, and Shuai Cheng Li. Supertad-fast: Accelerating topologically associating domains detection through discretization. *Journal of Computational Biology*, 31(9), 2024.
- [28] Daliang Liu, Kun Yue, Wenjie Liu, Xiang Chen, and Liang Duan. Structural entropy based spatio-temporal sequence forecasting. In *Proceedings of the International Conference on Database Systems for Advanced Applications*, pages 693–702, 2025.
- [29] Daliang Liu, Kun Yue, Wenjie Liu, Xiang Chen, Liang Duan, Angsheng Li, and Philip S. Yu. Hierarchical text classification optimization via structural entropy and singular smoothing. *IEEE Transactions on Knowledge and Data Engineering*, 37(9):5283–5297, 2025.
- [30] Qitong Liu, Hao Peng, Zuchen Li, Xihang Meng, Ziyu Yang, Jiting Li, Li Sun, and Philip S. Yu. Effective and unsupervised social event detection and evolution via RAG and structural entropy, 2026.
- [31] Xuecheng Liu, Luoyi Fu, Xinbing Wang, and Chenghu Zhou. On the similarity between von neumann graph entropy and structural information: Interpretation, computation, and applications. *IEEE Transactions on Information Theory*, 2022.
- [32] Yiwei Liu, Jiamou Liu, Zijian Zhang, and et al. Rem: From structural entropy to community structure deception. In *Proceedings of the NeurIPS*, 2019.
- [33] Mark EJ Newman and Michelle Girvan. Finding and evaluating community structure in networks. *Physical review E*, 69(2):026113, 2004.
- [34] Andrew Y. Ng, Michael I. Jordan, and Yair Weiss. On spectral clustering: analysis and an algorithm. In *Advances in Neural Information Processing Systems*, 2002.
- [35] OpenAI. Openai o3-mini: Pushing the frontier of cost-effective reasoning, 2025.
- [36] Lawrence Page. The pagerank citation ranking: Bringing order to the web. Technical report, Stanford InfoLab, 1999.
- [37] Yicheng Pan, Bingchen Fan, Pengyu Long, and Feng Zheng. An information-theoretic perspective of hierarchical clustering on graphs. In *Proceedings of the Forty-first Conference on Uncertainty in Artificial Intelligence*, volume 286 of *Proceedings of Machine Learning Research*, pages 3322–3345, 2025.
- [38] Yicheng Pan, Feng Zheng, and Bingchen Fan. An information-theoretic perspective of hierarchical clustering. *arXiv preprint arXiv:2108.06036*, 2021.

- [39] Judah Pearl. *Causality: models, reasoning and inference*. Cambridge University Press, 2nd edition, 2009.
- [40] Cheng Peng, Yonghao Li, Wanfu Gao, Jie Wen, and Weiping Ding. Sehfs: Structural entropy-guided high-order correlation learning for multi-view multi-label feature selection, 2026.
- [41] Hao Peng, Jingyun Zhang, Xiang Huang, and et al. Unsupervised social bot detection via structural information theory. *ACM TOIS*, 2024.
- [42] Kun Peng, Cong Cao, Hao Peng, Zhifeng Hao, Lei Jiang, Kongjing Gu, Yanbing Liu, and Philip S. Yu. Dialogues aspect-based sentiment quadruple extraction via structural entropy minimization partitioning, 2025.
- [43] Pascal Pons and Matthieu Latapy. Computing communities in large networks using random walks. In *Computer and Information Sciences-ISCIS*, 2005.
- [44] Pascal Pons and Matthieu Latapy. Computing communities in large networks using random walks. *Journal of Graph Algorithms and Applications*, 2006.
- [45] Yuyang Ren, Haonan Zhang, Luoyi Fu, and et al. Hi-part: Going beyond graph pooling with hierarchical partition tree for graph-level representation learning. *ACM TKDD*, 2024.
- [46] Jorma Rissanen. Modeling by shortest data description. *Automatica*, 1978.
- [47] Martin Rosvall and Carl T Bergstrom. Maps of random walks on complex networks reveal community structure. *Proceedings of the national academy of sciences*, 2008.
- [48] Claude E Shannon. A mathematical theory of communication. *The Bell System Technical Journal*, 27(3):379–423, 1948.
- [49] Jianbo Shi and Jitendra Malik. Normalized cuts and image segmentation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 22(8):888–905, 2000.
- [50] Kaiyu Song, Kun Yue, Liang Duan, Mingze Yang, and Angsheng Li. Mutual information based bayesian graph neural network for few-shot learning. In *Proceedings of the UAI*. PMLR, 2022.
- [51] Dingli Su, Hao Peng, Yicheng Pan, and Angsheng Li. A survey of structural entropy: Theory, methods, and applications. In *Proceedings of the 34th International Joint Conference on Artificial Intelligence (IJCAI)*, 2025. Camera-ready version, June 2025.
- [52] Dingli Su, Hao Peng, Guangjie Zeng, Pu Li, Angsheng Li, and Yicheng Pan. Emergence of cooperation in multi-agent reinforcement learning via coalition labeling and structural entropy. In *Proceedings of the SDM*, 2025.
- [53] Li Sun, Zhenhao Huang, and et al. Lsenet: Lorentz structural entropy neural network for deep graph clustering. In *Proceedings of the ICML*, 2024.
- [54] Li Sun, Zhenhao Huang, Yujie Wang, Hongbo Lv, Chunyang Liu, Hao Peng, and Philip S. Yu. Isosel: Isometric structural entropy learning for deep graph clustering in hyperbolic space, 2025.

- [55] Yongkai Tian, Xin Yu, Yirong Qi, Li Wang, Pu Feng, Wenjun Wu, and Jie Luo. Structural entropy guided hierarchical symmetric multi-agent reinforcement learning. *Expert Systems with Applications*, page 132289, 2026.
- [56] Naftali Tishby, Fernando C Pereira, and William Bialek. The information bottleneck method. In *Proceedings of the 37th annual Allerton conference on communication, control, and computing*, pages 368–377, 2000.
- [57] Vincent A. Traag, Ludo Waltman, and Nees Jan van Eck. From louvain to leiden: guaranteeing well-connected communities. *Scientific Reports*, 9:5233, 2019.
- [58] Anton Tsitsulin, John Palowitch, Bryan Perozzi, and Emmanuel Müller. Graph clustering with graph neural networks. In *Proceedings of the International Conference on Learning Representations*, 2020.
- [59] Chenxu Wang, Chaozhuo Li, Pengbo Wang, Litian Zhang, Songyang Liu, Ji Qi, Jiahui Hu, Yushan Cai, Hao Zhao, and Rui Pu. Lancet: Neural intervention via structural entropy for mitigating faithfulness hallucinations in LLMs, 2026.
- [60] Jiahui Wang, Jiangming Liu, Kun Yue, Liang Duan, and Jiande Wu. Structural entropy guided relation extraction on adaptive graph structure. *Pattern Recognition*, 172:112789, 2026.
- [61] Linqin Wang, Yaping Liu, Zhengtao Yu, and et al. Secodec: Structural entropy-based compressive speech representation codec for speech language models. In *Proceedings of the AAAI*, 2025.
- [62] Linqin Wang, Zhengtao Yu, Shengxiang Gao, and et al. Voice conversion via structural entropy. In *Proceedings of the IEEE ICASSP*, 2025.
- [63] Xinbing Wang, Huquan Kang, Luoyi Fu, and et al. Quantifying knowledge from the perspective of information structurization. *Plos one*, 2023.
- [64] Yifei Wang, Yupan Wang, Zeyu Zhang, Song Yang, Kaiqi Zhao, and Jiamou Liu. User: Unsupervised structural entropy-based robust graph neural network. In *Proceedings of the AAAI*, volume 37, pages 10235–10243, 2023.
- [65] Junran Wu, Xueyuan Chen, Bowen Shi, Shangzhe Li, and Ke Xu. Sega: Structural entropy guided anchor view for graph contrastive learning. In *ICML*, 2023.
- [66] Junran Wu, Xueyuan Chen, Ke Xu, and Shangzhe Li. Structural entropy guided graph hierarchical pooling. In *Proceedings of the ICML*, 2022.
- [67] Yantuan Xian, Pu Li, and et al. Community detection in large-scale complex networks via structural entropy game. In *Proceedings of the Web Conference*, 2025.
- [68] Minhui Xie, Hao Peng, Pu Li, and et al. Hierarchical superpixel segmentation via structural information theory. In *Proceedings of the SDM*, 2025.
- [69] Tianchi Xie, Jiangning Zhu, Guozu Ma, Minzhi Lin, Wei Chen, Weikai Yang, and Shixia Liu. Structural-entropy-based sample selection for efficient and effective learning. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [70] Runze Yang, Hao Peng, Angsheng Li, and et al. Hierarchical abstracting graph kernel. *IEEE TKDE*, 2024.

- [71] Runze Yang, Hao Peng, Chunyang Liu, and Angsheng Li. Incremental measurement of structural entropy for dynamic graphs. *Artificial Intelligence*, 2024.
- [72] Yingguang Yang, Qi Wu, Buyun He, and et al. Sebot: Structural entropy guided multi-view contrastive learning for social bot detection. In *Proceedings of the ACM SIGKDD*, 2024.
- [73] Zhenyu Yang, Ge Zhang, Jia Wu, and et al. Minimum entropy principle guided graph neural networks. In *Proceedings of the ACM WSDM*, 2023.
- [74] Zhiwei Yang, Haimei Qin, Xiaoyan Yu, Hao Peng, Lei Jiang, Li Sun, and Zhiqin Yang. Structural entropy guided incremental learning for open-world multimodal social event detection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 16119–16127, 2026.
- [75] Zhiwei Yang, Yuecen Wei, Haoran Li, and et al. Adaptive differentially private structural entropy minimization for unsupervised social event detection. In *Proceedings of the ACM CIKM*, 2024.
- [76] Yuan Yao, Yicheng Pan, Shaojiang Wang, Hongan Wang, and Waqas Nazeer. Flow-based clustering on directed graphs: A structural entropy minimization approach. *IEEE Access*, 2019.
- [77] Zhenqiang Ye, Jinjie Lu, Tianlong Gu, Fengrui Hao, and Xuemin Wang. Fairgse: Fairness-aware graph neural network without high false positive rates. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pages 16163–16171, 2026.
- [78] Zhitao Ying, Jiaxuan You, Christopher Morris, Xiang Ren, William L. Hamilton, and Jure Leskovec. Hierarchical graph representation learning with differentiable pooling. In *Advances in Neural Information Processing Systems*, 2018.
- [79] Xiaoyan Yu, Yifan Wei, Shuaishuai Zhou, and et al. Towards effective, efficient and unsupervised social event detection in the hyperbolic space. In *Proceedings of the AAAI*, 2025.
- [80] Guangjie Zeng, Hao Peng, and et al. Semi-supervised clustering via structural entropy with different constraints. In *Proceedings of the SDM*, 2024.
- [81] Guangjie Zeng, Hao Peng, Angsheng Li, and et al. Unsupervised skin lesion segmentation via structural entropy minimization on multi-scale superpixel graphs. In *Proceedings of the ICDM*, 2023.
- [82] Guangjie Zeng, Hao Peng, Angsheng Li, and et al. Scalable semi-supervised clustering via structural entropy with different constraints. *IEEE TKDE*, 2025.
- [83] Xianghua Zeng, Hao Peng, and Angsheng Li. Effective and stable role-based multi-agent collaboration by structural information principles. In *Proceedings of the AAAI*, 2023.
- [84] Xianghua Zeng, Hao Peng, and Angsheng Li. Adversarial socialbots modeling based on structural information principles. In *Proceedings of the AAAI*, 2024.
- [85] Xianghua Zeng, Hao Peng, and Angsheng Li. Effective exploration based on the structural information principles. In *Proceedings of the NeurIPS*, 2024.

- [86] Xianghua Zeng, Hao Peng, and Angsheng Li. Robustness evaluation of graph-based news detection using network structural information. In *Proceedings of the ACM SIGKDD*, 2025.
- [87] Xianghua Zeng, Hao Peng, and Angsheng Li. Proactive bot detection based on structural information principles. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 48(4):4726–4743, 2026.
- [88] Xianghua Zeng, Hao Peng, Angsheng Li, and et al. Hierarchical state abstraction based on structural information principles. In *Proceedings of the IJCAI*, 2023.
- [89] Xianghua Zeng, Hao Peng, Yicheng Pan, and Angsheng Li. Hyperbolic continuous structural entropy for hierarchical clustering. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 28094–28102, 2026.
- [90] Xianghua Zeng, Hao Peng, Yicheng Pan, Angsheng Li, and Guanlin Wu. Structural information-based hierarchical diffusion for offline reinforcement learning. In *Advances in Neural Information Processing Systems*, 2025.
- [91] Xianghua Zeng, Hao Peng, DingLi Su, and Angsheng Li. Hierarchical decision making based on structural information principles. *JMLR*, 2025.
- [92] Chong Zhang, Cheng Deng, Luoyi Fu, and et al. Rwe: A random walk based graph entropy for the structural complexity of directed networks. *IEEE TNSE*, 2024.
- [93] Chong Zhang, Qiang Guo, Luoyi Fu, and et al. Finding the source in networks: an approach based on structural entropy. *ACMT TOIT*, 2023.
- [94] Jingyun Zhang, Li Sun, Wu, Zhengtao Yu, and et al. Unsupervised graph clustering with deep structural entropy. In *Proceedings of the ACM SIGKDD*, 2025.
- [95] Jun Zhang, Liang Duan, Xiang Chen, Xinhui Li, Jing Chai, and Kun Yue. Structural entropy based graph contrastive learning for node classification. *Adaptive Behavior*, 2025.
- [96] Siyu Zhang, Jiamou Liu, Yiwei Liu, Zijian Zhang, and Bakhadyr Khoussainov. Improving togetherness using structural entropy. In *Proceedings of the MobiCASE*. Springer, 2021.
- [97] Yu Wei Zhang, Meng Bo Wang, and Shuai Cheng Li. Supertad: robust detection of hierarchical topologically associated domains with optimized structural information. *Genome biology*, 2021.
- [98] Zaijun Zhang, Daoyun Xu, and Jincheng Zhou. A structural entropy measurement principle of propositional formulas in conjunctive normal form. *Entropy*, 2021.
- [99] Zehan Zhang, Jiahuan Yan, Jie Zhang, Junjie Liu, Yuanchen Bei, Shijie Zhang, Xiangnan He, and Tat-Seng Chua. Enhanced pre-training for recommendation via hypergraph structural entropy. *ACM Transactions on Information Systems*, 2025.
- [100] Tianle Zhao, Yida Zhang, Zhiyuan Xie, Tianchi Xie, Wei Chen, Yalong Yang, Jing Zhang, and Shixia Liu. Sese: Black-box uncertainty quantification for large language models based on structural information theory, 2025.

- [101] He Zhu, Chong Zhang, Junjie Huang, Junran Wu, and Ke Xu. Hitin: Hierarchy-aware tree isomorphism network for hierarchical text classification. In *Proceedings of the ACL*, 2023.
- [102] Dongcheng Zou, Hao Peng, Xiang Huang, and et al. Se-gsl: A general and effective graph structure learning framework through structural entropy optimization. In *Proceedings of the ACM Web Conference*, 2023.
- [103] Dongcheng Zou, Senzhang Wang, Xuefeng Li, Hao Peng, Yuandong Wang, Chunyang Liu, Kehua Sheng, and Bo Zhang. Multispans: A multi-range spatial-temporal transformer network for traffic forecast via structural entropy optimization. In *Proceedings of the ACM WSDM*, 2024.
- [104] Weihui Zou, Xianghua Zeng, Jianxin Li, Hao Peng, and Angsheng Li. A structural information guided hierarchical reconstruction for graph anomaly detection. In *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management*, pages 3685–3694, 2024.

A Survey Protocol and Companion Library

A.1 Literature Search Protocol

This survey covers SE-related publications from January 2016 (the publication date of the foundational SE paper [16]) through May 2026. We conducted systematic searches on Google Scholar, Semantic Scholar, arXiv, and the ACL/ACM/IEEE digital libraries using the following query terms: “structural entropy”, “encoding tree”, “SE minimization”, “hierarchical structural entropy”, and “Li Angsheng”. We included papers that (a) directly use SE as a primary methodology, (b) propose extensions of the SE formalism, or (c) introduce SE-based applications. Papers using the term “structural entropy” in a purely figurative sense without connection to the Li–Pan framework [16] were excluded.

A.2 The `glass-jax` Companion Library

`glass-jax` (Graph-based Latent Abstraction & Structural Segmentation in JAX) is a fully differentiable, hardware-accelerated framework for graph clustering that serves as this survey’s reproducibility artifact.

Core objectives.

- **Soft Modularity:** relaxes the Louvain objective via $Q = \frac{1}{2m} \text{Tr}(S^\top BS)$ with soft assignment matrix S .
- **Soft Map Equation:** a differentiable approximation of Infomap computing the description length of a random walk over soft assignments.
- **Soft 2D Structural Entropy (Glass-SE):** soft-assignment implementation of Definition 3.2, enabling gradient flow through the SE objective.

Optimization engine. The *multi-start vmap engine* runs B parallel optimization trajectories using `jax.vmap`, initialized with 1 spectral embedding and $B-1$ random Gaussian initializations. Temperature annealing decays τ from 1.0 to 0.01 over 500 Adam steps,

sharpening soft assignments into near-discrete decisions. The trajectory achieving the best (unscaled) objective value is selected.

Benchmarking interface. All experiments in Section 6 are fully reproducible via:

```
git clone https://github.com/SELGroup/glass-jax
cd glass-jax
pip install -e ".[dev]"
python experiments/benchmark_same_protocol.py
```

Requirements. $\text{JAX} \geq 0.4$, `optax`, `networkx`, `scikit-learn`, `python-louvain`, `infomap` (C++ wrapper). GPU/TPU acceleration is automatic via JAX's XLA compiler.

Table 3: Cross-objective evaluation (mean $\pm$ 95% CI over 5 seeds). Modularity $\uparrow$; Map Equation and Structural Entropy $\downarrow$. Bold: best per column per dataset.

Dataset	Algorithm	Modularity $\uparrow$	Map Eq. $\downarrow$	Struct. Ent. $\downarrow$
Karate	Louvain	0.440$\pm$0.006	4.165 $\pm$ 0.044	3.323$\pm$0.015
	Infomap	0.435 $\pm$ 0.000	4.087$\pm$0.000	3.495 $\pm$ 0.000
	SEP	0.404 $\pm$ 0.000	4.198 $\pm$ 0.000	3.731 $\pm$ 0.000
	LSENet	0.404 $\pm$ 0.000	4.198 $\pm$ 0.000	3.731 $\pm$ 0.000
	Glass-Mod	0.404 $\pm$ 0.000	4.198 $\pm$ 0.000	3.731 $\pm$ 0.000
	Glass-Map	0.224 $\pm$ 0.083	4.680 $\pm$ 0.193	3.983 $\pm$ 0.135
	Glass-SE	0.404 $\pm$ 0.000	4.198 $\pm$ 0.000	3.731 $\pm$ 0.000
Caveman	Louvain	0.878$\pm$0.000	3.548$\pm$0.000	3.394$\pm$0.000
	Infomap	0.878$\pm$0.000	3.548$\pm$0.000	3.394$\pm$0.000
	SEP	0.878$\pm$0.000	3.548$\pm$0.000	3.394$\pm$0.000
	LSENet	0.870 $\pm$ 0.008	3.623 $\pm$ 0.080	3.471 $\pm$ 0.083
	Glass-Mod	0.852 $\pm$ 0.008	3.807 $\pm$ 0.075	3.664 $\pm$ 0.081
	Glass-Map	0.824 $\pm$ 0.000	3.972 $\pm$ 0.000	3.827 $\pm$ 0.000
	Glass-SE	0.860 $\pm$ 0.010	3.729 $\pm$ 0.098	3.585 $\pm$ 0.106
SBM (Clean)	Louvain	0.418$\pm$0.006	6.909$\pm$0.022	6.011$\pm$0.009
	Infomap	0.418$\pm$0.006	6.909$\pm$0.022	6.011$\pm$0.009
	SEP	0.408 $\pm$ 0.010	6.945 $\pm$ 0.037	6.026 $\pm$ 0.015
	LSENet	0.418$\pm$0.006	6.909$\pm$0.022	6.011$\pm$0.009
	Glass-Mod	0.418$\pm$0.006	6.909$\pm$0.022	6.011$\pm$0.009
	Glass-Map	0.317 $\pm$ 0.072	7.112 $\pm$ 0.149	6.275 $\pm$ 0.191
	Glass-SE	0.418$\pm$0.006	6.909$\pm$0.022	6.011$\pm$0.009
SBM (Noisy)	Louvain	0.212$\pm$0.007	7.758 $\pm$ 0.045	6.202$\pm$0.005
	Infomap	0.000 $\pm$ 0.000	7.187$\pm$0.004	7.187 $\pm$ 0.004
	SEP	0.167 $\pm$ 0.007	7.757 $\pm$ 0.023	6.400 $\pm$ 0.015
	LSENet	0.204 $\pm$ 0.009	7.652 $\pm$ 0.031	6.337 $\pm$ 0.016
	Glass-Mod	0.205 $\pm$ 0.011	7.644 $\pm$ 0.036	6.337 $\pm$ 0.020
	Glass-Map	0.128 $\pm$ 0.038	7.562 $\pm$ 0.079	6.623 $\pm$ 0.133
	Glass-SE	0.208 $\pm$ 0.011	7.638 $\pm$ 0.036	6.330$\pm$0.018

Table 4: Clustering on real-world citation networks (ACC, NMI, ARI). †: topology-only method; ‡: feature-augmented method. Best per metric per dataset in bold.

Dataset	Algorithm	ACC	NMI	ARI
Cora ($N=2,708$, $K=7$)	Louvain [†]	0.372	0.439	0.236
	Glass-SE [†]	0.274	0.076	0.039
	LSENet [‡]	0.387	0.266	0.164
Citeseer ($N=3,327$, $K=6$)	Louvain [†]	0.192	0.329	0.094
	Glass-SE [†]	0.252	0.037	0.025
	LSENet [‡]	0.403	0.195	0.166